

WHEN ARE TRIANGLES INVISIBLE? FUNDAMENTAL LIMITS OF HIGHER-ORDER STRUCTURE IDENTIFIABILITY FROM EDGE FLOWS

Weiping Yan

University of Minnesota, Twin Cities
Minneapolis, MN, USA

ABSTRACT

Topological signal processing places signals on the triangles of a simplicial complex, but which triangles are “filled” is often unknown and must be inferred from edge flows. Existing work gives algorithms; we ask when the task is *possible at all*. The latent structure is observed only through the *curl* of the flows, which annihilates gradient and harmonic nuisances; classical detection theory then yields a curl-invisibility threshold $\rho^*(N) = \Theta(\sqrt{\log(1/\delta)/N})$, a $\log \binom{p}{k}$ Fano converse for joint recovery, and exact per-triangle marginal laws with effective SNR $\sigma_c^2/(\sigma_n^2(\mathbf{G}^{-1})_{\tau\tau})$ after a GLS-type decorrelation. Our main result corrects the folklore “rank obstruction”: it is a *first- vs. second-order identifiability dichotomy*. Equal-curl-subspace supports are indistinguishable under deterministic signals — on K_n the curl subspace carries a $3/n$ degrees-of-freedom ratio — yet under random signals *every* support is identifiable at *any* rank deficiency, because the lifted signatures $\{\mathbf{b}_\tau, \mathbf{b}_\tau^\top\}$ are always linearly independent. Unfavorable geometry costs sample complexity, not possibility: minimal confusers (tetrahedral face swaps) have universal curl-domain separation 16, giving exponent $\rho_2^2(1 + o(1))$, $\rho_2 = \sigma_c^2/\sigma_n^2$. All constants are Monte-Carlo validated, and the theory recovers genuine, unplanted structure from real data — tropical-cyclone circulation from ERA5 wind-field edge flows, validated against full-resolution vorticity and the independent IBTrACS archive. Code and data are released.

Index Terms— Topological signal processing, simplicial complexes, Hodge decomposition, identifiability, detection theory, edge flows.

1. INTRODUCTION

Graph signal processing captures pairwise structure, but many networked measurements—currency exchange, traffic circulation, neural flow, statistical ranking—carry *higher-order* structure that lives on the triangles of a simplicial complex [1, 2, 3]. Topological signal processing (TSP) has rapidly built a toolbox on the Hodge Laplacian: convolutional filters [4], maximally localized dictionaries [5], Dirac operators [6], adaptive estimation [7], and unrolled reconstruction [8]. A recurring primitive is that the set of filled triangles is unknown and must be learned from edge-flow signals.

Prior work. Learning a graph topology from node signals is well charted [9, 10]. Its higher-order analogue — inferring which triangles are filled — has been treated *algorithmically*: MAP estimation [11], greedy topology learning [12], sparse cell-complex augmentation [13]; adaptive schemes [7] list structure discovery as open. HodgeRank [14] established the curl as the signature of higher-order inconsistency. Missing is the *converse side*: when

is the task information-theoretically solvable, at what snapshot budget, and how does the answer depend on the geometry? (Simplicial/hypergraph SBM detectability observes the structure *directly*; here it is latent, seen only through the flows it shapes.) We supply that converse side. Our tools are deliberately classical — detection theory and Fano arguments [15, 16], in the lineage of support-recovery and sparse-covariance detection limits [17, 18, 19, 20] — our contribution is the reduction that makes them bite, plus one dichotomy theorem (Thm. 3) we believe is genuinely new:

1. A *curl-annihilation* reduction (Lemma 1; the detection-theoretic use of HodgeRank’s observation): the latent structure is observable only through the curl, giving a standard two-variance test (Prop. 1) and an *exact* finite-sample product law over edge-disjoint candidates (Prop. 2).
2. Converse thresholds through the reduction (Thm. 1): the sample-complexity scaling $\rho^*(N) = \Theta(\sqrt{\log(1/\delta)/N})$ at the exponent level, and a Fano $\log \binom{p}{k}$ factor for joint recovery.
3. *Geometry-aware decorrelation* for edge-sharing candidates — the GLS/BLUE inversion $\hat{\mathbf{y}} = \mathbf{G}^{-1}\mathbf{c}$ [21] — with exact marginal laws whose effective SNR is the variance-inflation factor’s (Thm. 2), plus a union-bound guarantee (Cor. 1).
4. Main theorem: a *first- vs. second-order identifiability dichotomy* (Thm. 3) correcting the folklore rank obstruction: equal-image supports are indistinguishable under deterministic signals (on K_n : a $3/n$ DoF ratio), yet under random signals *every* support is identifiable at *any* rank deficiency (lifted-spark Lemma 2), at a price with universal constants (swap separation 16; exponent $\rho_2^2(1 + o(1))$).

Every result is Monte-Carlo validated. On real data we recover *genuine, unplanted* structure — cyclone circulation from ERA5 edge flows, validated against full-resolution vorticity and the independent IBTrACS archive — with FX as a machine-precision consistency check and road networks grounding the recovery laws on real geometry.

2. BACKGROUND AND SIGNAL MODEL

Fix a graph $(\mathcal{V}, \mathcal{E})$ with node–edge incidence $\mathbf{B}_1 \in \mathbb{R}^{|\mathcal{V}| \times |\mathcal{E}|}$ and, for a candidate triangle set $\mathcal{T}$ (e.g. all 3-cliques), edge–triangle incidence $\mathbf{B}_2 \in \mathbb{R}^{|\mathcal{E}| \times |\mathcal{T}|}$; N is the snapshot count, $p = |\mathcal{T}|$. The identity $\mathbf{B}_1\mathbf{B}_2 = \mathbf{0}$ holds, and the Hodge 1-Laplacian $\mathbf{L}_1 = \mathbf{B}_1^\top\mathbf{B}_1 + \mathbf{B}_2\mathbf{B}_2^\top$ induces the orthogonal split

$$\mathbb{R}^{|\mathcal{E}|} = \underbrace{\text{im } \mathbf{B}_1^\top}_{\text{gradient}} \oplus \underbrace{\text{im } \mathbf{B}_2}_{\text{curl}} \oplus \underbrace{\ker \mathbf{L}_1}_{\text{harmonic}}. \quad (1)$$




Fig. 1. Curl-invisibility recovery threshold. Colour is empirical $P(\text{exact recovery})$; the solid line is the exact finite-sample law (Prop. 2), the dotted line the Bonferroni-corrected Chernoff floor of Thm. 1 ($\delta \rightarrow \delta/p$).




Fig. 2. Edge-sharing strip. Left: mean Hamming error vs. N ; the naive detector never recovers. Right: decorrelated detector vs. the exact marginal theory — independence approximation (dashed) and rigorous union bound (dotted).

The *curl* of a flow is $\text{curl } \mathbf{f} = \mathbf{B}_2^\top \mathbf{f}$. Each column of $\mathbf{B}_2$ has three ± 1 entries, so $\mathbf{G} = \mathbf{B}_2^\top \mathbf{B}_2$ has $G_{\tau\tau} = 3$, $G_{\sigma\tau} \in \{0, \pm 1\}$, nonzero iff σ, τ share an edge. The 1-skeleton is known; the latent object is the filled set $\mathcal{S} \subseteq \mathcal{T}$. We observe N i.i.d. snapshots

$$\mathbf{f}_t = \mathbf{B}_1^\top \mathbf{a}_t + \mathbf{B}_{2,\mathcal{S}} \mathbf{y}_t + \mathbf{h}_t + \mathbf{n}_t, \quad t = 1, \dots, N, \quad (2)$$

with $\mathbf{a}_t \sim \mathcal{N}(\mathbf{0}, \sigma_g^2 \mathbf{I})$ (gradient nuisance), $\mathbf{y}_t \sim \mathcal{N}(\mathbf{0}, \sigma_c^2 \mathbf{I})$ on $\mathcal{S}$, harmonic $\mathbf{h}_t \in \ker \mathbf{L}_1$, noise $\mathbf{n}_t \sim \mathcal{N}(\mathbf{0}, \sigma_n^2 \mathbf{I})$; the goal is to recover $\mathcal{S}$. This is the generative family of the algorithmic literature (cf. [11]); we characterize when its inverse problem is solvable.

3. FUNDAMENTAL LIMITS

3.1. Curl annihilation and the single-triangle test

Lemma 1 (Curl annihilation). *For any $\mathbf{f}$, the statistic $\mathbf{B}_2^\top \mathbf{f}$ depends only on the curl component: $\mathbf{B}_2^\top (\mathbf{B}_1^\top \mathbf{a}) = \mathbf{0}$ for all $\mathbf{a}$ and $\mathbf{B}_2^\top \mathbf{h} = \mathbf{0}$ for all $\mathbf{h} \in \ker \mathbf{L}_1$. Under (2),*

$$\mathbf{c}_t = \mathbf{B}_2^\top \mathbf{f}_t = (\mathbf{B}_2^\top \mathbf{B}_{2,\mathcal{S}}) \mathbf{y}_t + \mathbf{B}_2^\top \mathbf{n}_t. \quad (3)$$

Proof. $\mathbf{B}_1 \mathbf{B}_2 = \mathbf{0}$ gives $\mathbf{B}_2^\top \mathbf{B}_1^\top = \mathbf{0}$. If $\mathbf{L}_1 \mathbf{h} = \mathbf{0}$ then $0 = \mathbf{h}^\top \mathbf{L}_1 \mathbf{h} \geq \mathbf{h}^\top \mathbf{B}_2 \mathbf{B}_2^\top \mathbf{h} = \|\mathbf{B}_2^\top \mathbf{h}\|^2$, so $\mathbf{B}_2^\top \mathbf{h} = \mathbf{0}$. $\square$

Thus the nuisances—however large—are removed exactly; all information about $\mathcal{S}$ flows through (3). For an *isolated* candidate τ

(no shared edges) the scalar $c_{\tau,t}$ is zero-mean Gaussian with variance $v_0 = 3\sigma_n^2$ if $\tau \notin \mathcal{S}$ and $v_1 = 9\sigma_c^2 + 3\sigma_n^2 = v_0(1 + \rho)$ if $\tau \in \mathcal{S}$, where

$$\rho \triangleq 3\sigma_c^2/\sigma_n^2 \quad (4)$$

is the *curl signal-to-noise ratio*.

Proposition 1 (Two-variance detection). *Detecting an isolated τ from N snapshots is a two-variance Gaussian test: the energy $E = \sum_{t=1}^N c_{\tau,t}^2$ is sufficient, $E/v_i \sim \chi_N^2$ under H_i , the minimum Bayes error is attained by thresholding E and admits a closed form via χ_N^2 tails, and it decays as $\exp(-N C(v_0, v_1))$, where $C(v_0, v_1) = \max_{s \in [0,1]} g(s)$ is the Gaussian Chernoff information [15, 16], with $w_i = 1/v_i$ and*

$$g(s) = \frac{1}{2} [\log(s w_0 + (1-s) w_1) - s \log w_0 - (1-s) \log w_1]. \quad (5)$$

Theorem 1 (Curl-invisibility threshold; joint-recovery converse). *(a) With ρ as in (4), $C(\rho) = \rho^2/16 + o(\rho^2)$ as $\rho \rightarrow 0$. Hence for target error δ and budget N , attaining the target error exponent requires $\rho \geq \rho^*(N) = \Theta(\sqrt{\log(1/\delta)/N})$ for any estimator. (b) Recovering a uniformly random size- k support among p edge-disjoint candidates with error $\leq \epsilon$ requires $N \geq ((1-\epsilon) \log \binom{p}{k} - \log 2)/(k \text{KL}_1)$ with $\text{KL}_1 = (\rho - \log(1+\rho))/2 \leq \rho^2/4$; joint recovery pays a $\log \binom{p}{k}$ factor over the single-triangle budget.*

Proof sketch. (a) Chernoff information is the optimal Bayes error exponent for i.i.d. observations [16]; Taylor expansion gives $\rho^2/16$, and $e^{-NC(\rho)} \leq \delta$ forces the scaling. (b) Independent curl coordinates give $I(\mathcal{S}; \mathbf{c}^N) \leq Nk \text{KL}(\mathcal{N}(0, v_1) \parallel \mathcal{N}(0, v_0))$; Fano finishes (supplement §S1, incl. a signal-agnostic variant). Honest scope: (a) is exponent-level — the finite- N 50% boundary sits an order-one factor below the floor (same $1/\sqrt{N}$ law, Fig. 1); the scaling is the classical variance-detection rate [15, 17]; specific to our problem are the reduction producing ρ and the constants (3, 9). $\square$

When the candidates are pairwise edge-disjoint, the coordinates of $\mathbf{c}_t$ are *independent* across triangles, so support recovery decouples exactly:

Proposition 2 (Exact finite-sample recovery law). *For edge-disjoint candidates and the per-triangle Bayes-optimal energy threshold, $P(\hat{\mathcal{S}} = \mathcal{S}) = \prod_{\tau \in \mathcal{S}} (1 - P_{\text{miss}}) \prod_{\tau \notin \mathcal{S}} (1 - P_{\text{fa}})$, with $P_{\text{miss}}, P_{\text{fa}}$ given by χ_N^2 tails at the common threshold. Its 50% contour in the (N, ρ) plane is the empirical recovery boundary (Fig. 1).*

3.2. Confusable triangles: geometry-aware decorrelation

Edge-sharing candidates are *confusable*: an active triangle leaks curl energy onto neighbours through the off-diagonal entries of $\mathbf{G}$, so naive curl-energy thresholding fails (Fig. 2, left) and *worsens* as ρ grows. The repair is the textbook GLS/BLUE inversion of the curl map [21] (for invertible $\mathbf{G}$, OLS, GLS, BLUE and ML all coincide at $\mathbf{G}^{-1} \mathbf{c}$); what matters here is the *exact finite-sample marginal detection law* it induces, with a geometry factor that is precisely the classical variance-inflation factor.

Theorem 2 (Decorrelated marginal laws). *Assume $\mathbf{B}_2$ has full column rank, so $\mathbf{G}$ is invertible. Let $\hat{\mathbf{y}}_t = \mathbf{G}^{-1} \mathbf{c}_t$ and let $\tilde{\mathbf{y}}_t$ denote $\mathbf{y}_t$ zero-padded to all of $\mathcal{T}$. Then, exactly,*

$$\hat{\mathbf{y}}_t = \tilde{\mathbf{y}}_t + \mathbf{e}_t, \quad \mathbf{e}_t \sim \mathcal{N}(\mathbf{0}, \sigma_n^2 \mathbf{G}^{-1}). \quad (6)$$

Hence the marginal law of each coordinate $\hat{y}_{\tau,t}$ is a two-variance test with $v_{0,\tau} = \sigma_n^2(\mathbf{G}^{-1})_{\tau\tau}$, $v_{1,\tau} = \sigma_c^2 + v_{0,\tau}$, and per-triangle effective SNR

$$\rho_\tau^{\text{eff}} = \frac{\sigma_c^2}{\sigma_n^2(\mathbf{G}^{-1})_{\tau\tau}}. \quad (7)$$

Proof. By Lemma 1, $\mathbf{c}_t = \mathbf{G}\tilde{\mathbf{y}}_t + \mathbf{B}_2^\top \mathbf{n}_t$ (the columns of $\mathbf{G}$ indexed by $\mathcal{S}$ act on $\mathbf{y}_t$; zero-padding absorbs the indexing). Applying $\mathbf{G}^{-1}$ gives the mean exactly; the noise covariance is $\sigma_n^2 \mathbf{G}^{-1} \mathbf{G} \mathbf{G}^{-1} = \sigma_n^2 \mathbf{G}^{-1}$. Gaussian marginals of (6) give the per-coordinate two-variance laws. $\square$

$\mathbf{G}^{-1}$ decorrelates the *mean* (removing leakage exactly); the noise stays correlated — $\sigma_n^2 \mathbf{G}^{-1}$ is not diagonal — hence no “whitening”, and the joint law does *not* factorize. $(\mathbf{G}^{-1})_{\tau\tau}$ is the price of geometry: $1/3$ when isolated (recovering $\rho_\tau^{\text{eff}} = \rho$), growing as signatures overlap. What survives arbitrary correlations is a union bound over exact marginals:

Corollary 1 (Rigorous recovery guarantee). *Threshold each $\sum_t \hat{y}_{\tau,t}^2$ at its marginal-optimal (or Bonferroni) level and let $P_{\text{err},\tau}$ be the resulting marginal error probabilities from Theorem 2. Then $\Pr[\hat{\mathcal{S}} = \mathcal{S}] \geq 1 - \sum_{\tau \in \mathcal{T}} P_{\text{err},\tau}$, regardless of the correlation structure of $\mathbf{e}_t$.*

Remark 1 (Independence approximation). *The product $\prod_\tau (1 - P_{\text{err},\tau})$ is exact when the scores are independent (diagonal $\mathbf{G}$; Prop. 2) and an approximation otherwise; on our diagonally-dominant benchmarks it tracks Monte-Carlo within 0.03 across the transition, while the union bound stays a valid, more conservative floor (Fig. 2, right).*

3.3. The rank obstruction is a dichotomy, not impossibility

When curl signatures are linearly dependent ($\mathbf{G}$ singular — e.g. all $\binom{n}{3}$ triangles of K_n), the folklore answer — “identifiable only modulo $\ker \mathbf{B}_2$ ” — is correct for *deterministic* signals but *false* for the random model (2): the covariance carries $\sigma_c^2 \sum_{\tau \in \mathcal{S}} \mathbf{b}_\tau \mathbf{b}_\tau^\top$, strictly finer than the column space. The correct statement is a dichotomy, resting on a spark property of the *lifted* signatures:

Lemma 2 (Lifted spark). *For any candidate set of distinct 3-cliques of a simple graph, the matrices $\{\mathbf{b}_\tau \mathbf{b}_\tau^\top\}_{\tau \in \mathcal{T}}$ are linearly independent.*

Proof. A pair of distinct edges lies in at most one common triangle, so for $e \neq e'$ both in τ , the (e, e') entry of $\sum_\tau w_\tau \mathbf{b}_\tau \mathbf{b}_\tau^\top$ equals $\pm w_\tau$: every coefficient is read off a distinct matrix entry. $\square$

Theorem 3 (First- vs. second-order identifiability). *Let $r = \text{rank } \mathbf{B}_2$ and write $\rho_2 = \sigma_c^2 / \sigma_n^2$. (i) First order. If $\text{im } \mathbf{B}_{2,\mathcal{S}} = \text{im } \mathbf{B}_{2,\mathcal{S}'}$, then under unknown deterministic triangle signals the two supports induce identical families of flow distributions — indistinguishable at any SNR and any N . Such confusers exist whenever the candidate set contains all four faces of a tetrahedron ($\mathbf{b}_{012} - \mathbf{b}_{013} + \mathbf{b}_{023} - \mathbf{b}_{123} = \mathbf{0}$); on K_n , $\dim \text{im } \mathbf{B}_2 = \binom{n-1}{2}$, a $3/n$ DoF ratio. (ii) Second order. Under the random model (2), the map $\mathcal{S} \mapsto \text{Cov}(\mathbf{f})$ is injective (Lemma 2), so every support is identifiable from i.i.d. snapshots — at any rank deficiency, including all of K_n . (iii) The price. For a minimal confuser pair (a tetrahedral face swap $\mathcal{S}' = \mathcal{S} \setminus \{\tau_a\} \cup \{\tau_b\}$), the optimal test between the two curl-coordinate covariances has Chernoff exponent $C_G = \rho_2^2 (1 + o(1))$ as $\rho_2 \rightarrow 0$, since the curl-domain separation is universal: $\|\mathbf{u}_a \mathbf{u}_a^\top - \mathbf{u}_b \mathbf{u}_b^\top\|_F^2 = 9 + 9 - 2(\mathbf{u}_a^\top \mathbf{u}_b)^2 = 16$ (faces*

share exactly one edge), and the Bhattacharyya expansion gives $C_G = (\rho_2^2/16) \cdot 16(1 + o(1))$. Hence distinguishing the pair costs $N^ \sim \log(1/\delta)/\rho_2^2$ snapshots, and a Fano argument over the $4\binom{n}{4}$ leave-one-face-out hypotheses adds a $\log n$ factor for full support recovery on K_n .*

Proof sketch. (i) Equal images make the deterministic mean families coincide. (ii) In curl coordinates $\mathbf{z}_t = \mathbf{Q}^\top \mathbf{f}_t$ ($\mathbf{Q}$ orthonormal basis of $\text{im } \mathbf{B}_2$) the law is $\mathcal{N}(\mathbf{0}, \sigma_n^2 \mathbf{I}_r + \sigma_c^2 \sum_{\tau \in \mathcal{S}} \mathbf{u}_\tau \mathbf{u}_\tau^\top)$, $\mathbf{u}_\tau = \mathbf{Q}^\top \mathbf{b}_\tau$; injectivity transfers from Lemma 2 since $\mathbf{b}_\tau \in \text{im } \mathbf{B}_2$. (iii) Direct computation plus the second-order (Bhattacharyya) expansion of the Chernoff information; constants Monte-Carlo-validated in the released tests ($C_G/\rho_2^2 \rightarrow 1$ to 0.998; separation 16 exact, exhaustively minimal over all equal-image pairs on K_5). Full proofs: supplement §S4. $\square$

Theorem 3 reconciles theory with practice: second-order estimators — fitting the empirical curl covariance with the atoms $\{\mathbf{g}_\tau \mathbf{g}_\tau^\top\}$, greedily or by non-negative lasso — are the achievers of regime (ii), verified on rank-deficient K_5 . The $3/n$ ratio survives as first-order observability; unfavorable geometry is a quantifiable snapshot premium, not a wall.

Remark 2 (Open). *Sharp first-order sparse-support thresholds over $\ker \mathbf{B}_2$ (spark/coherence conditions) and the exact constant in (iii)’s $\log n$ factor remain open.*

4. EXPERIMENTS

All experiments run on CPU from fixed seeds; every theoretical curve is overlaid on independent Monte-Carlo simulation, and a released test-suite checks each closed form against simulation. Code and data reproduce all figures via a single script.

4.1. Recovery threshold (validation of Thm. 1, Prop. 2)

Fig. 1 sweeps (N, ρ) on an edge-disjoint planted complex ($p = 8$, 4 active, gradient/harmonic nuisance present). Recovery transitions rapidly (the steepness is the finite- p product of sigmoids — no critical-point claim); the exact product law of Prop. 2 tracks the empirical 50% contour to a median ratio 1.01; the dotted Bonferroni-corrected Chernoff floor ($\delta \rightarrow \delta/p$) shows the $1/\sqrt{N}$ law.

4.2. Confusability (validation of Thm. 2, Cor. 1)

On an edge-sharing strip ($p = 9$, alternating 4 active, $\rho = 8$; full column rank), the naive energy detector plateaus at Hamming error ≈ 2.6 regardless of N , while the decorrelated detector of Theorem 2 recovers exactly (Fig. 2, left); its recovery probability tracks the independence approximation within Monte-Carlo error and respects the union bound, conservative in the transition as a bound should be (right). A greedy residual-fitting baseline (standing in for [12]) stalls at ≈ 0.3 .

4.3. Second-order identifiability where subspaces fail (Thm. 3)

On K_5 ($p = 10$, $\text{rank } \mathbf{B}_2 = 6$: where a subspace argument forbids recovery), the greedy covariance-atom estimator recovers a planted 4-triangle support exactly ($N = 1200$, gradient nuisance present); exhaustively, all equal-image support pairs of K_4 have distinct covariances and every tetrahedral swap has separation exactly 16. The Monte-Carlo error of the optimal confuser-pair test respects the $e^{-N C_G}$ envelope with empirical exponent converging onto C_G , and $C_G/\rho_2^2 \in [0.95, 1.0]$ for $\rho_2 \leq 0.01$ (released tests).




Fig. 3. Unplanted recovery of real higher-order structure. (A) Decorrelated curl scores on the mesh for one 4-day window; red circles: independent IBTrACS cyclone fixes. (B) ROC of the (fully plug-in) detector against both ground truths across all windows. (C) Detection quality vs. snapshot budget N , with the theory floor $\rho^*(N)$ overlaid.




Fig. 4. Real road-network geometry. (A) Full-column-rank B_2 : DoF ratio 1 (green) vs. K_9 's $3/n$ (red). (B) Anaheim ($G = 3I$): planted recovery follows the exact edge-disjoint product law (df $N-1$ after centering out the real background). (C) EMA (non-diagonal G): the geometry-aware detector's heterogeneous marginal laws at work on real geometry — empirical vs. independence approximation and union bound.

4.4. Real, unplanted recovery: cyclones from wind-field edge flows

The achievability side is exercised on *genuine* latent structure (Fig. 3): tropical-cyclone circulation in ERA5 winds [22], Western North Pacific, Aug–Sep 2020. Ten-meter winds (0.7° grid) become edge flows on a triangulated mesh (line integrals); by Stokes' theorem a triangle's curl is its circulation — area-integrated vorticity — so high-curl triangles are real atmospheric vortices. The mesh is simply connected, so B_2 has full column rank by Euler's formula: no confusers, the favorable regime of Theorem 3. Per 4-day window (16 six-hourly snapshots) the temporally centered curl-energy statistic of Prop. 1, area-normalized to the vorticity scale, scores each triangle; *nothing is planted, no oracle parameters* (the decorrelated variant is reported in the repository: its G^{-1} noise amplification hurts pure detection ranking on this near-regular mesh). Validation is against two independent references: full-resolution finite-difference vorticity (never seen by the mesh detector) and the IBTrACS best-track archive [23], agency-verified positions independent of the reanalysis (panel B). Degrading the snapshot budget traces the predicted $1/\sqrt{N}$ scaling (panel C).

4.5. Geometry on real infrastructure; FX as a consistency check

Fig. 4 grounds the recovery laws on real road geometry (TNTP [24]). (A) All three networks have full-column-rank B_2 (DoF ratio 1 vs. $3/n$ on K_9); the real equilibrium flows carry genuine rotational energy (2.4%/4.6% curl). (B) Anaheim is edge-disjoint, so $G = 3I$ and planted recovery *collapses to the product law* (Prop. 2); the real constant background is removed exactly by centering (df $N-1$) and

the nuisances are annihilated, so — stated plainly — the surviving statistical problem is synthetic; the honest real recovery evidence is Fig. 3. (C) EMA exercises Theorem 2 on real geometry: 10 edge-sharing pairs make G non-diagonal; empirical recovery tracks the independence approximation and respects the union bound. *FX consistency check (no figure)*: edge flows from 257 days of ECB rates on K_9 are by construction a single-potential gradient $f = B_1^\top p$, so Lemma 1 predicts *identically zero* curl; measured curl/gradient $= 1.8 \times 10^{-31}$, floating-point residue. A consistency check on real market data (HodgeRank's arbitrage-free reading [14]), not a market discovery; K_9 's 28/84 rank is panel (A)'s red bar.

5. CONCLUSION

We characterized when higher-order structure is identifiable from edge flows: a curl-invisibility threshold with a $\log \binom{p}{k}$ joint converse; exact marginal laws after GLS-type decorrelation with a union-bound guarantee; and a first- vs. second-order dichotomy — rank deficiency binds only for deterministic signals; random signals make every support identifiable at a quantified snapshot premium. This delimits when algorithmic TSP inference [11, 12, 13] succeeds, with achievability shown on genuine, unplanted structure (ERA5 cyclone circulation, two independent references). The repository adds a signal-agnostic Fano variant, sub-Gaussian achievability, an exact q^{3k} edge-sampling law [25], and plug-in estimation. Remark 2's thresholds remain open.

6. REFERENCES

- [1] Sergio Barbarossa and Stefania Sardellitti, “Topological signal processing over simplicial complexes,” *IEEE Transactions on Signal Processing*, vol. 68, pp. 2992–3007, 2020.
- [2] Michael T. Schaub, Yu Zhu, Jean-Baptiste Seby, T. Mitchell Roddenberry, and Santiago Segarra, “Signal processing on higher-order networks: Livin’ on the edge... and beyond,” *Signal Processing*, vol. 187, pp. 108149, 2021.
- [3] Lek-Heng Lim, “Hodge Laplacians on graphs,” *SIAM Review*, vol. 62, no. 3, pp. 685–715, 2020.
- [4] Maosheng Yang, Elvin Isufi, Michael T. Schaub, and Geert Leus, “Simplicial convolutional filters,” *IEEE Transactions on Signal Processing*, vol. 70, pp. 4633–4648, 2022.
- [5] Claudio Battiloro, Paolo Di Lorenzo, and Sergio Barbarossa, “Topological Slepians: Maximally localized representations of signals over simplicial complexes,” in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2023.
- [6] Lucille Calmon, Michael T. Schaub, and Ginestra Bianconi, “Dirac signal processing of higher-order topological signals,” *New Journal of Physics*, vol. 25, pp. 093013, 2023.
- [7] Lorenzo Marinucci, Claudio Battiloro, and Paolo Di Lorenzo, “Topological adaptive least mean squares algorithms over simplicial complexes,” *IEEE Transactions on Signal and Information Processing over Networks*, vol. 12, pp. 283–298, 2026, arXiv:2505.23160.
- [8] Chengen Liu, Geert Leus, and Elvin Isufi, “Unrolling of simplicial Elasticnet for edge flow signal reconstruction,” *IEEE Open Journal of Signal Processing*, vol. 5, 2024, DOI 10.1109/OJSP.2023.3339376.
- [9] Xiaowen Dong, Dorina Thanou, Michael Rabbat, and Pascal Frossard, “Learning graphs from data: A signal representation perspective,” *IEEE Signal Processing Magazine*, vol. 36, no. 3, pp. 44–63, 2019.
- [10] Gonzalo Mateos, Santiago Segarra, Antonio G. Marques, and Alejandro Ribeiro, “Connecting the dots: Identifying network structure via graph signal processing,” *IEEE Signal Processing Magazine*, vol. 36, no. 3, pp. 16–43, 2019.
- [11] Sravanthi Gurugubelli and Sundeep Prabhakar Chepuri, “Simplicial complex learning from edge flows via sparse clique sampling,” in *European Signal Processing Conference (EUSIPCO)*, 2024.
- [12] Andrei Buciulea, Elvin Isufi, Geert Leus, and Antonio G. Marques, “Learning the topology of a simplicial complex using simplicial signals: A greedy approach,” arXiv:2502.20159, 2025.
- [13] Josef Hoppe and Michael T. Schaub, “Representing edge flows on graphs via sparse cell complexes,” arXiv:2309.01632; Learning on Graphs Conference (LoG), 2023.
- [14] Xiaoye Jiang, Lek-Heng Lim, Yuan Yao, and Yinyu Ye, “Statistical ranking and combinatorial Hodge theory,” *Mathematical Programming*, vol. 127, no. 1, pp. 203–244, 2011.
- [15] H. Vincent Poor, *An Introduction to Signal Detection and Estimation*, Springer, 2nd edition, 1994.
- [16] Thomas M. Cover and Joy A. Thomas, *Elements of Information Theory*, Wiley, 2nd edition, 2006.
- [17] Martin J. Wainwright, “Information-theoretic limits on sparsity recovery in the high-dimensional and noisy setting,” *IEEE Transactions on Information Theory*, vol. 55, no. 12, pp. 5728–5741, 2009.
- [18] Arash A. Amini and Martin J. Wainwright, “High-dimensional analysis of semidefinite relaxations for sparse principal components,” *The Annals of Statistics*, vol. 37, no. 5B, pp. 2877–2921, 2009.
- [19] Quentin Berthet and Philippe Rigollet, “Optimal detection of sparse principal components in high dimension,” *The Annals of Statistics*, vol. 41, no. 4, pp. 1780–1815, 2013.
- [20] T. Tony Cai, Cun-Hui Zhang, and Harrison H. Zhou, “Optimal rates of convergence for covariance matrix estimation,” *The Annals of Statistics*, vol. 38, no. 4, pp. 2118–2144, 2010.
- [21] Steven M. Kay, *Fundamentals of Statistical Signal Processing: Estimation Theory*, Prentice-Hall, 1993.
- [22] Hans Hersbach, Bill Bell, Paul Berrisford, et al., “The ERA5 global reanalysis,” *Quarterly Journal of the Royal Meteorological Society*, vol. 146, no. 730, pp. 1999–2049, 2020.
- [23] Kenneth R. Knapp, Michael C. Kruk, David H. Levinson, Howard J. Diamond, and Charles J. Neumann, “The International Best Track Archive for Climate Stewardship (IBTrACS): Unifying tropical cyclone data,” *Bulletin of the American Meteorological Society*, vol. 91, no. 3, pp. 363–376, 2010.
- [24] Transportation Networks for Research Core Team, “Transportation networks for research,” github.com/bstabler/TransportationNetworks, 2023, Sioux Falls, Eastern-Massachusetts, and Anaheim networks with user-equilibrium flow solutions.
- [25] Siheng Chen, Rohan Varma, Aliaksei Sandryhaila, and Jelena Kovačević, “Discrete signal processing on graphs: Sampling theory,” *IEEE Transactions on Signal Processing*, vol. 63, no. 24, pp. 6510–6523, 2015.